

# rnn\_music\_gen

May 1, 2019

```
[1]: import tensorflow as tf
    tf.enable_eager_execution()

    import numpy as np
    import os
    import time
    import matplotlib.pyplot as plt

[2]: # Read, then decode for py2 compat.
    text = open('data.txt', 'rb').read().decode(encoding='utf-8')
    # length of text is the number of characters in it
    print ('Length of text: {} characters'.format(len(text)))
```

Length of text: 2072753 characters

```
[3]: # The unique characters in the file
    vocab = sorted(set(text))
    print ('{} unique characters'.format(len(vocab)))
```

34 unique characters

```
[4]: # Creating a mapping from unique characters to indices
    char2idx = {u:i for i, u in enumerate(vocab)}
    idx2char = np.array(vocab)

    text_as_int = np.array([char2idx[c] for c in text])

[5]: # The maximum length sentence we want for a single input in characters
    seq_length = 100
    examples_per_epoch = len(text)//seq_length

    # Create training examples / targets
    char_dataset = tf.data.Dataset.from_tensor_slices(text_as_int)
    sequences = char_dataset.batch(seq_length+1, drop_remainder = True)

    for item in sequences.take(5):
        print(repr(''.join(idx2char[item.numpy()])))
```

```
'<music21.note.Note B->\n<music21.chord.Chord D4 F3 B-3>\n<music21.chord.Chord
D5 F5 B-5 D6>\n<music21.ch
'ord.Chord F3 D4 B-3>\n<music21.chord.Chord F5 B-5 D5 D6>\n<music21.chord.Chord
D4 B-3 F3>\n<music21.chor
'd.Chord D5 E-6 F5 B-5>\n<music21.chord.Chord D4 B-3 F3>\n<music21.chord.Chord
D6 D5 B-5 F5>\n<music21.ch
'ord.Chord B-3 D4 F3>\n<music21.chord.Chord F5 D5 B-5 D6>\n<music21.chord.Chord
D4 F3 B-3>\n<music21.chor
'd.Chord D5 B-4 F5 B-5>\n<music21.chord.Chord F3 D3 B-3>\n<music21.note.Note
B->\n<music21.chord.Chord B-
```

```
[6]: def split_input_target(chunk):
      input_text = chunk[:-1]
      target_text = chunk[1:]
      return input_text, target_text

      dataset = sequences.map(split_input_target)

[7]: for input_example, target_example in dataset.take(1):
      print ('Input data: ', repr(''.join(idx2char[input_example.numpy()])))
      print ('Target data:', repr(''.join(idx2char[target_example.numpy()])))
```

```
Input data:  '<music21.note.Note B->\n<music21.chord.Chord D4 F3
B-3>\n<music21.chord.Chord D5 F5 B-5 D6>\n<music21.c'
Target data: 'music21.note.Note B->\n<music21.chord.Chord D4 F3
B-3>\n<music21.chord.Chord D5 F5 B-5 D6>\n<music21.ch'
```

```
[8]: for i, (input_idx, target_idx) in enumerate(zip(input_example[:5],
→target_example[:5])):
      print("Step {:4d}".format(i))
      print("  input: {} ({:s})".format(input_idx, repr(idx2char[input_idx])))
      print(" expected output: {} ({:s})".format(target_idx,
→repr(idx2char[target_idx])))
```

```
Step    0
  input: 12 ('<')
 expected output: 27 ('m')
Step    1
  input: 27 ('m')
 expected output: 33 ('u')
Step    2
  input: 33 ('u')
 expected output: 31 ('s')
Step    3
  input: 31 ('s')
 expected output: 26 ('i')
Step    4
```

```
input: 26 ('i')
expected output: 22 ('c')
```

```
[9]: # Batch size
BATCH_SIZE = 256
steps_per_epoch = examples_per_epoch//BATCH_SIZE

# Buffer size to shuffle the dataset
# (TF data is designed to work with possibly infinite sequences,
# so it doesn't attempt to shuffle the entire sequence in memory. Instead,
# it maintains a buffer in which it shuffles elements).
BUFFER_SIZE = 10000

dataset = dataset.shuffle(BUFFER_SIZE).batch(BATCH_SIZE, drop_remainder=True)

dataset
```

```
[9]: <BatchDataset shapes: ((256, 100), (256, 100)), types: (tf.int64, tf.int64)>
```

```
[10]: # Length of the vocabulary in chars
vocab_size = len(vocab)

# The embedding dimension
embedding_dim = 256

# Number of RNN units
rnn_units = 1024
```

```
[11]: model = tf.keras.Sequential()
model.add(tf.keras.layers.Embedding(len(vocab), embedding_dim,
                                     batch_input_shape=[BATCH_SIZE, None]))
model.add(tf.keras.layers.CuDNNGRU(rnn_units,
                                     return_sequences=True,
                                     recurrent_initializer='glorot_uniform',
                                     stateful=True))
model.add(tf.keras.layers.Dense(len(vocab)))

print(model.summary())
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(256, None, 256)	8704
cu_dnngru (CuDNNGRU)	(256, None, 1024)	3938304
dense (Dense)	(256, None, 34)	34850
Total params: 3,981,858		

Trainable params: 3,981,858

Non-trainable params: 0

-----  
None

```
[12]: for input_example_batch, target_example_batch in dataset.take(1):  
      example_batch_predictions = model(input_example_batch)  
      print(example_batch_predictions.shape, "# (batch_size, sequence_length, vocab_size)")
```

(256, 100, 34) # (batch\_size, sequence\_length, vocab\_size)

```
[13]: def loss(labels, logits):  
      return tf.nn.sparse_softmax_cross_entropy_with_logits(labels=labels, logits=logits)  
  
model.compile(  
    optimizer = tf.train.AdamOptimizer(),  
    loss = loss)
```

```
[14]: # Directory where the checkpoints will be saved  
checkpoint_dir = './training_music_checkpoints'  
# Name of the checkpoint files  
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt_c_{epoch}")  
  
checkpoint_callback=tf.keras.callbacks.ModelCheckpoint(  
    filepath=checkpoint_prefix,  
    save_weights_only=True)
```

```
[15]: EPOCHS=10  
  
# history = model.fit(dataset.repeat(), epochs=EPOCHS, steps_per_epoch=steps_per_epoch, callbacks=[checkpoint_callback])  
  
history = model.fit(dataset.repeat(), epochs=EPOCHS, steps_per_epoch=steps_per_epoch, callbacks=[checkpoint_callback])
```

Epoch 1/10

80/80 [=====] - 14s 170ms/step - loss: 1.3489

Epoch 2/10

80/80 [=====] - 14s 171ms/step - loss: 0.2433

Epoch 3/10

80/80 [=====] - 14s 170ms/step - loss: 0.2139

Epoch 4/10

80/80 [=====] - 14s 171ms/step - loss: 0.2028

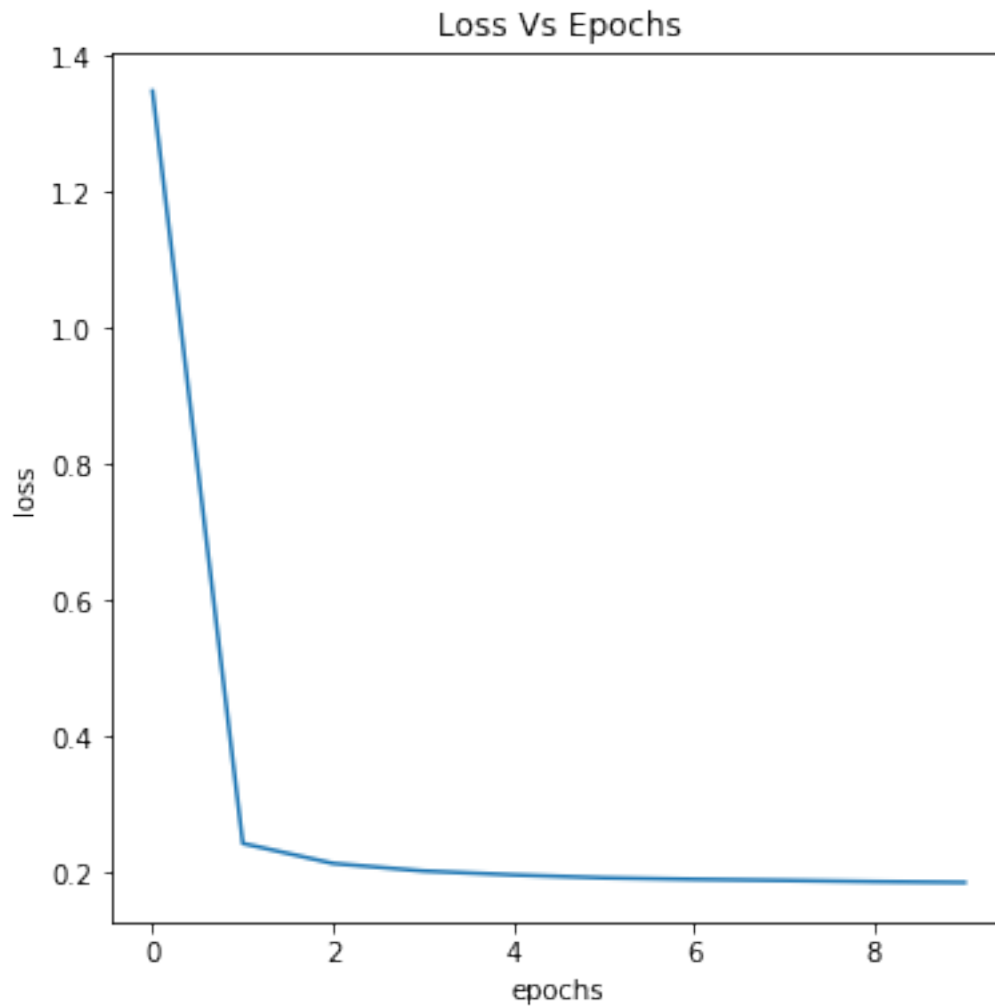
Epoch 5/10

80/80 [=====] - 14s 173ms/step - loss: 0.1971

Epoch 6/10

```
80/80 [=====] - 14s 172ms/step - loss: 0.1928
Epoch 7/10
80/80 [=====] - 14s 172ms/step - loss: 0.1905
Epoch 8/10
80/80 [=====] - 14s 174ms/step - loss: 0.1889
Epoch 9/10
80/80 [=====] - 14s 175ms/step - loss: 0.1872
Epoch 10/10
80/80 [=====] - 14s 175ms/step - loss: 0.1860
```

```
[16]: plt.figure(figsize = (6,6))
plt.plot(history.history['loss'])
plt.title('Loss Vs Epochs')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.show()
```



```
[18]: checkpoint_dir = './training_music_checkpoints/ckpt_c_10'
      # tf.train.latest_checkpoint(checkpoint_dir)

      model = tf.keras.Sequential()
      model.add(tf.keras.layers.Embedding(len(vocab), embedding_dim,
                                          batch_input_shape=[1, None]))
      model.add(tf.keras.layers.CuDNNGRU(rnn_units,
                                          return_sequences=True,
                                          recurrent_initializer='glorot_uniform',
                                          stateful=True))
      model.add(tf.keras.layers.Dense(len(vocab)))

      model.load_weights(checkpoint_dir)

      model.build(tf.TensorShape([1, None]))

      print(model.summary())
```

```
-----
Layer (type)                 Output Shape              Param #
=====
embedding_2 (Embedding)      (1, None, 256)           8704
-----
cu_dnngru_2 (CuDNNGRU)       (1, None, 1024)          3938304
-----
dense_2 (Dense)              (1, None, 34)            34850
=====
Total params: 3,981,858
Trainable params: 3,981,858
Non-trainable params: 0
-----
None
```

```
[21]: def generate_text(model, start_string):
      # Evaluation step (generating text using the learned model)

      # Number of characters to generate
      num_generate = 5000

      # Converting our start string to numbers (vectorizing)
      input_eval = [char2idx[s] for s in start_string]
      input_eval = tf.expand_dims(input_eval, 0)

      # Empty string to store our results
      text_generated = []

      # Low temperatures results in more predictable text.
```

```

# Higher temperatures results in more surprising text.
# Experiment to find the best setting.
temperature = 1.0

# Here batch size == 1
model.reset_states()
for i in range(num_generate):
    predictions = model(input_eval)
    # remove the batch dimension
    predictions = tf.squeeze(predictions, 0)

    # using a multinomial distribution to predict the word returned by the
    →model
    predictions = predictions / temperature
    predicted_id = tf.multinomial(predictions, num_samples=1)[-1,0].numpy()

    # We pass the predicted word as the next input to the model
    # along with the previous hidden state
    input_eval = tf.expand_dims([predicted_id], 0)

    text_generated.append(idx2char[predicted_id])

return (start_string + ''.join(text_generated))

```

```

[23]: with open('gen.txt', 'w') as f:
        f.write(generate_text(model, start_string=u"<"))

```